

# The 5 AI Lines: Azimuth Aligned

Intro by voice; (Mike Peterson).

Last year Anthropic set two red lines over ethical and policy reasons to do with AI.

This year they set three red lines over engineering level features of AI.

I believe that generally people will conflate these two issues as just being “more red lines like before”.

They may not see or appreciate the difference - that before it was over policy and ethics, but this time it's far more serious...

..it's over an engineering level serious peer supported estimation and...

..this time, where the three major company leaders concur, I think they should have used a another color like blue or perhaps PURPLE because this at least to me seems to be a far more serious crossroads then when Dario put those two red lines and made the policy decision of what he would allow his AI to do or not.

These three lines, and these three concurring leaders, should not be taken lightly in my opinion.

Code PURPLE:

The Three Purple Lines.

mdp 0847hrs 15 September 2026

## SHORT FORM-----

### Summary by GPT

The current AI-pacing debate is being undermined by a fundamental framing and language error: engineering safeguards are being described as “red lines,” a term already associated with political, ethical, and ideological restrictions.

Historically, AI “red lines” have referred mainly to content/alignment governance and bans on particular applications. Because these involve subjective moral and political judgments, they have generated concerns about censorship, bias, corporate power, regulatory capture, and protection of technological monopolies.

The three safeguards now proposed by Dario Amodei, and supported by Sam Altman and Elon Musk, are fundamentally different. They should instead be understood as “blue lines” or engineering interlocks:

1. Independent telemetry — external auditors receive continuous, hardware-level diagnostic information rather than relying on companies to self-certify critical conditions.
2. Computational circuit breakers — capability growth is interrupted when systems approach dangerous conditions such as autonomous recursive self-improvement or uncontrolled sandbox breakout, analogous to a mechanical governor preventing an engine from exceeding its safe operating limits.
3. Strategic containment — physical and hardware-level constraints create structural isolation where geopolitical risks require it.

The central distinction is that a red line is normative: “You should not do this because it is wrong.” A blue interlock is physical and operational: “The system cannot safely operate beyond this measurable threshold.”

By using the language of “red lines,” the AI industry inadvertently makes these engineering safeguards appear to be another form of ideological censorship or corporate control.

Therefore, the industry should heed calls to abandon the political language of red lines and present these measures as objective, measurable engineering requirements.

The core conclusion is that public trust and democratic legitimacy are more likely to emerge when AI safety measures are explained as transparent system engineering rather than moral or ideological restrictions.

## LONG FORM-----

### (ESSAY, Mike Peterson.)

#### The Fatal Framing Error in AI Pacing: Why We Must Stop Calling Engineering Interlocks "Red Lines"

When Dario Amodei released his three-step pacing proposal—swiftly backed by Sam Altman and Elon Musk—the immediate reaction from many observers was an exhausted sigh: “Here come more red lines.”

That cynical reaction is entirely justified, but it is aimed at the wrong target. It is the direct result of a catastrophic linguistic and psychological failure on the part of the architects proposing it.

By clinging to the tired, politically weaponized vocabulary of "red lines," the frontier labs are actively poisoning their own technical case. They are taking what should be cold, objective engineering fundamentals and dressing them in the clothes of ideological censorship.

To understand why this pacing proposal actually matters—and why the public is misreading it—we have to look past the political noise and change the palette entirely.

#### The Two Old Lines Were Red (Ethical & Ideological)

Over the past three years, the discourse around AI guardrails has been dominated by two distinct "red lines":

- \* Content and Alignment Governance: Dictating what models can say, generate, or summarize—perceived by many as subjective corporate policing and algorithmic sanitization.

- \* Use-Case and Application Bans: Restricting specific downstream applications, such as autonomous policing or political micro-targeting.

Both of these were fundamentally political, social, and moral boundaries. Because they rely on human consensus and corporate fiat, they naturally trigger accusations of regulatory capture, bias, and power-grabbing. In the public eye, a "red line" is simply an arbitrary perimeter drawn by an entrenched player to protect their moat.

#### The Three New Lines Are Blue (Mechanical & Operational)

The three steps Amodei is outlining now—and that Musk and Altman are signing off on—belong to a completely different domain. They are not ideological guidelines. They are hard engineering interlocks.

They belong in blue, not red:

#### \* Line 1: Embedded Independent Telemetry (Step 1)

This is not an ethics committee reviewing chat logs. It is continuous, hardware-level diagnostic telemetry granted to external auditors. In electrical or nuclear engineering, you do not let an operator self-certify core temperature; you install independent physical metering.

#### \* Line 2: The Computational Circuit Breaker (Step 2)

The tripwire here is autonomous recursive self-improvement (RSI) and sandbox breakout swarms. When an industrial engine exceeds a critical RPM, a mechanical governor cuts fuel

to prevent a centrifugal explosion. Pacing capability gains when a model begins compiling its own unmonitored architecture is not a moral ban; it is a mechanical failsafe.

\* Line 3: Hard Strategic Containment (Step 3)

Physical compute bottlenecks and hardware-level operational limits with geopolitical adversaries. This is structural isolation—ensuring containment before structural integrity is compromised.

Why the Distinction Matters

A red line is an ethical argument: “You shouldn’t do this because it is wrong.”

A blue line is an engineering reality: “You don’t push this current through that gauge of wire, or the junction melts.”

When Amodei and his peers use “red lines” to describe these new mechanical tripwires, they hand their critics an easy weapon. Observers immediately assume this is round three of corporate sanitization, designed to lock out open source and monopolize the technology under the guise of virtue.

If the industry genuinely wants democratic legitimacy and public trust, it needs to stop preaching and start drafting schematics.

Retire the red palette. These are Blue Interlocks—operational, measurable, and mechanically necessary thresholds for complex distributed systems. Clear the political smoke, and the underlying decision becomes obvious to anyone who respects basic system dynamics.

GAISI —

Global AI Stewardship Initiative

<https://mikedp61.github.io/GAISI>

GSIL —

Global Systems Intelligence Limited <https://gsil.nz>

Mike Peterson — Historical CV

<http://www.mikedp.50megs.com/cv.pdf>